

Score-based Generative Modeling in Latent Space

Pramook Khungurn

June 14, 2023

This note is written as I read the paper “Score-based Generative Modeling in Latent Space” by Vahdat et al. [VKK21].

1 Introduction

- This paper presents an interesting approach to generative modeling. It is a combination of a variational autoencoder (VAE) and a denoising diffusion generative model (DDPM) or a score-based model (SGM).
- The paper starts with a VAE that is already quite good, the NVAE [VK21], and try to make it better by combining it with a diffusion model.
- For a VAE, the latent code \mathbf{z} is ideally supposed to be distributed according to a prior distribution $p(\mathbf{z})$, which we typically take to be the Gaussian distribution $\mathcal{N}(\mathbf{0}, I)$.
- Nevertheless, the distribution of latent codes that the VAE decoder can decode well is the distribution $\{\mathbf{z} : \mathbf{x} \sim p_{\text{data}}, \mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})\}$ where q_{ϕ} is the VAE encoder. This distribution might not agree with $\mathcal{N}(\mathbf{0}, I)$ even though we have optimized the VAE as best as we can.
- The idea of the paper is to transform $\mathcal{N}(\mathbf{0}, I)$ into $\{\mathbf{z} : \mathbf{x} \sim p_{\text{data}}, \mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})\}$ with a diffusion model.
- The approach is called “latent score-based generative model” (LSGM) because, to generate a sample, we use a diffusion model (aka a score-based model) to sample a VAE latent code. Then, we use the VAE decoder to generate a real data item.
- In this write-up, I will not follow the paper that closely. This is because I don’t like the way it formulates diffusion model, which uses the SDE in Song et al.’s 2021 paper [SSDK⁺21]. I’d rather use the VDM formulation [KSPH21], which I’m more familiar with.

2 Background

- VAE and DDPM are similar.
 - To generate a data sample, they both sample a **latent code** from a prior distribution.
 - Then, they transform the latent code to a real data sample through some kind of process.
- We shall call the process that goes from a latent code to a real data sample the “backward process.” The other way around is the “forward process”.
- The backward process is denoted by the probability function p , and the forward process the probability function q .
- We denote a real data sample with \mathbf{x} . We assume that it is a member of \mathbb{R}^D , which we call the **ambient space**. A data item is distributed according to the data distribution p_{data} .
- We denote a latent code with \mathbf{z} . It is a member of the **latent space** \mathbb{R}^d .

2.1 Variational Autoencoder

- For a VAE, we have that $d < D$ in general.
- The backward process.
 - Sample \mathbf{z} according to the prior distribution $p(\mathbf{z})$, which we generally take to be $\mathcal{N}(\mathbf{0}, I)$.
 - Use the VAE decoder to sample \mathbf{x} given \mathbf{z} .
 - So,

$$\begin{aligned} p(\mathbf{x}, \mathbf{z}) &= p(\mathbf{x}|\mathbf{z})p(\mathbf{z}) \\ p(\mathbf{x}) &= \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z}) d\mathbf{z}. \end{aligned} \tag{1}$$

- The forward process.
 - Sample \mathbf{x} according to p_{data} .
 - * For convenience, we shall denote p_{data} with q , so $p_{\text{data}}(\mathbf{x}) = q(\mathbf{x})$.
 - Use the encoder to sample \mathbf{z} given \mathbf{x} .
 - Symbolically,

$$q(\mathbf{x}, \mathbf{z}) = q(\mathbf{z}|\mathbf{x})q(\mathbf{x}).$$

- We model the encoder with a neural network with parameters ϕ . Typically, this network has two functions $\boldsymbol{\mu}_\phi(\mathbf{x})$ and $\boldsymbol{\sigma}_\phi(\mathbf{x})$, and the distribution it models is given by

$$q_\phi(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\mathbf{z}; \boldsymbol{\mu}_\phi(\mathbf{x}), \text{diag}(\boldsymbol{\sigma}_\phi^2(\mathbf{x}))).$$

- We model the decoder with a neural network with parameters ψ . The network only has one function $\boldsymbol{\mu}_\psi(\mathbf{z})$. The distribution it models is given by

$$p_\psi(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_\psi(\mathbf{z}), \sigma_E^2 I)$$

where σ_E is a hyperparameter.

- To optimize the network, we seek to minimize the negative log-likelihood

$$\begin{aligned} & E_{\mathbf{x} \sim q(\mathbf{x})} [-\log p(\mathbf{x})] \\ &= -E_{\mathbf{x} \sim q(\mathbf{x})} \left[\log \int p_\psi(\mathbf{x}|\mathbf{z})p(\mathbf{z}) d\mathbf{z} \right] && \text{(Equation 1)} \\ &= -E_{\mathbf{x} \sim q(\mathbf{x})} \left[\log \int q_\phi(\mathbf{z}|\mathbf{x}) \frac{p_\psi(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} d\mathbf{z} \right] && \text{(importance sampling)} \\ &= -E_{\mathbf{x} \sim q(\mathbf{x})} \left[\log E_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[\frac{p_\psi(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} \right] \right] && \text{(integral to expectation)} \\ &\geq -E_{\mathbf{x} \sim q(\mathbf{x})} \left[E_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \left[\log \frac{p_\psi(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})} \right] \right] && \text{(Jensen's inequality)} \\ &= -E_{\mathbf{x} \sim q(\mathbf{x}), \mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\psi(\mathbf{x}|\mathbf{z}) - \log q_\phi(\mathbf{z}|\mathbf{x}) + \log p(\mathbf{z})]. \end{aligned}$$

2.2 Variational Diffusion Model

- We follow the treatment of Kingma et al. [KSPH21], which is summarized in one of my notes [Khu22].
- For a diffusion model, $D = d$.
- The forward process starts with a data item \mathbf{x} sampled from $p_{\text{data}}(\mathbf{x}) = q(\mathbf{x})$.
- We then scale \mathbf{x} down and add noise to it. This process is a continuous process that runs from time $t = 0$ to $t = 1$. At time t , the noised data item is denoted by \mathbf{z}_t . The distribution of \mathbf{z}_t is given by:

$$q(\mathbf{z}_t|\mathbf{x}) = \mathcal{N}(\mathbf{z}_t; \alpha_t \mathbf{x}, \sigma_t^2 I).$$

The functions α_t and σ_t are together called the **noise schedule** of the diffusion process. The process is governed by the following stochastic differential equation:

$$d\mathbf{z}_t = \frac{\alpha'_t}{\alpha_t} \mathbf{z}_t dt + \sqrt{2\sigma_t \left(\sigma'_t - \frac{\alpha'_t}{\alpha_t} \right)} d\mathbf{W}.$$

- We require that $\alpha_t^2 + \sigma_t^2 = 1$ for all $t \in [0, 1]$. So, we are using the variance preserving formulation.
- The **signal-to-noise ratio** at time t , denoted by $\text{SNR}(t)$, is defined as:

$$\text{SNR}(t) = \frac{\alpha_t^2}{\sigma_t^2}.$$

3 Method

- The paper refers to the method it proposes as the **latent score-based generative model** (LSGM).
- The method contains the following components.
 - An encoder $q_\phi(\mathbf{z}_0|\mathbf{x})$.
 - A diffusion model $p_\theta(\mathbf{z}_0)$, which acts as the prior distribution.
 - A decoder $p_\psi(\mathbf{x}|\mathbf{z}_0)$
- To train the whole system, we would like to minimize the log-likelihood:

$$\begin{aligned} & E_{\mathbf{x} \sim p_{\text{data}}} [-\log p(\mathbf{x})] \\ & \leq E_{\mathbf{x} \sim p_{\text{data}}} \left[E_{\mathbf{z}_0 \sim q_\phi(\mathbf{z}_0|\mathbf{x})} [-\log p_\psi(\mathbf{x}|\mathbf{z}_0)] + \text{KL}(q_\phi(\mathbf{z}_0|\mathbf{x}) \| p_\theta(\mathbf{z}_0)) \right] \\ & = E_{\mathbf{x} \sim p_{\text{data}}, \mathbf{z}_0 \sim q_\phi(\mathbf{z}_0|\mathbf{x})} \left[-\log p_\psi(\mathbf{x}|\mathbf{z}_0) + \log q_\phi(\mathbf{z}_0|\mathbf{x}) - \log p_\theta(\mathbf{z}_0) \right]. \end{aligned}$$

The second line comes from the standard evidence lower-bound derivation found in VAE papers [Khu20]:

$$-\log p(\mathbf{x}) \leq E_{\mathbf{z}_0 \sim q_\phi(\mathbf{z}_0|\mathbf{x})} [-\log p_\psi(\mathbf{x}|\mathbf{z}_0)] + \text{KL}(q_\phi(\mathbf{z}_0|\mathbf{x}) \| p_\theta(\mathbf{z}_0)).$$

The third line just expanding the KL divergence into its constituent parts:

$$\text{KL}(q_\phi(\mathbf{z}_0|\mathbf{x}) \| p_\theta(\mathbf{z}_0)) = E_{\mathbf{z}_0 \sim q_\phi(\mathbf{z}_0|\mathbf{x})} \left[\log \frac{q_\phi(\mathbf{z}_0|\mathbf{x})}{p_\theta(\mathbf{z}_0)} \right] = E_{\mathbf{z}_0 \sim q_\phi(\mathbf{z}_0|\mathbf{x})} \left[\log q_\phi(\mathbf{z}_0|\mathbf{x}) - \log p_\theta(\mathbf{z}_0) \right].$$

- Now, the paper went on at a great length to find an expression for the cross entropy term

$$E_{\mathbf{z}_0 \sim q_\phi(\mathbf{z}_0|\mathbf{x})} \left[-\log p_\theta(\mathbf{z}_0) \right].$$

However, we will refer to the derivation in the “Variational Diffusion Models” paper by Kingma et al. [KSPH21, Khu22] as a shortcut.

- To apply the shortcut, we need to reinterpret the notations.
 - We will use q to denote the “forward” process. The forward process now has three steps.
 - * Sample $\mathbf{x} \sim p_{\text{data}}$.
 - * Sample the latent code \mathbf{z}_0 corresponding to \mathbf{x} : $\mathbf{z}_0 \sim q(\mathbf{z}_0|\mathbf{x})$.
 - This process is modeled with the encoder q_ϕ .
 - * Run the diffusion process forward in time to obtain \mathbf{z}_1 from \mathbf{z}_0 : $\mathbf{z}_1 \sim q(\mathbf{z}_1|\mathbf{z}_0)$.
 - The parameters of the forward process is the noise schedule α_t and σ_t .
 - We assume that $\alpha_t^2 + \sigma_t^2 = 1$; i.e., the variance preserving formulation.
 - We do this so that we can use the derivation in the Kingma et al. paper.
 - Note, however, that the Vahdat et al. paper is not limited to the variance preserving formulation.
 - The backward process is denoted by p . It also has three steps.
 - * Sample \mathbf{z}_1 according to $p(\mathbf{z}_1)$, which is an isotropic Gaussian distribution.
 - * Run the diffusion process backward to obtain \mathbf{z}_0 from \mathbf{z}_1 : $\mathbf{z}_0 \sim p(\mathbf{z}_0|\mathbf{z}_1)$.
 - This part is modeled by the diffusion model p_θ .
 - * Sample a data item in the ambient space \mathbf{x} from the latent code \mathbf{z}_0 : $\mathbf{x} \sim p(\mathbf{x}|\mathbf{z}_0)$.
 - This is modeled by the decoder p_ψ .
- In practice, however, we will not run the diffusion process from exactly $t = 0$ to $t = 1$. Instead, we will run it from $t_{\min} \gtrapprox 0$ to $t_{\max} \lesssim 1$. The reason for this is so that, for any time in the interval $[t_{\min}, t_{\max}]$, the **signal-to-noise ratio** (SNR) $\frac{\alpha_t^2}{\sigma_t^2}$ is finite.
- So, in practice, the backward process is as follows.
 - Sample $\mathbf{z}_{t_{\max}}$ according to $p(\mathbf{z}_{t_{\max}})$, which we shall approximate with $\mathcal{N}(\mathbf{0}, \sigma_{t_{\max}}^2 I)$.
 - Run the diffusion process backward to obtain $\mathbf{z}_{t_{\min}}$ from $\mathbf{z}_{t_{\max}}$.
 - Use $\mathbf{z}_{t_{\min}}$ as an approximation for \mathbf{z}_0 , and then sample the data item in the ambient space $\mathbf{x} \sim p(\mathbf{x}|\mathbf{z}_{t_{\min}})$.

- Using the variational lower bound in [Khu22], we have that

$$-\log p(\mathbf{z}_0) \leq \text{KL}(q(\mathbf{z}_1|\mathbf{z}_0)||p(\mathbf{z}_{t_{\max}})) + E_{\mathbf{z}_{t_{\min}} \sim q(\mathbf{z}_{t_{\min}}|\mathbf{z}_0)} \left[-\log p(\mathbf{z}_0|\mathbf{z}_{t_{\min}}) \right] + \mathcal{L}_\infty(\mathbf{z}_0).$$

If we choose t_{\min} to be low enough and t_{\max} high enough, the first two terms should be negligible. So,

$$\begin{aligned} -\log p(\mathbf{z}_0) &\lesssim \mathcal{L}_\infty(\mathbf{z}_0) \\ &= \frac{1}{2} E_{\boldsymbol{\xi} \sim \mathcal{N}(\mathbf{0}, I)} \left[\int_{\text{SNR}_{\min}}^{\text{SNR}_{\max}} \left\| \mathbf{z}_0 - \hat{\mathbf{z}}_\theta(\underbrace{\alpha_{t(v)} \mathbf{z}_0 + \sigma_{t(v)} \boldsymbol{\xi}}_{\mathbf{z}_{t(v)}}, t(v)) \right\|^2 dv \right] \\ &= \frac{1}{2} E_{\boldsymbol{\xi} \sim \mathcal{N}(\mathbf{0}, I), v \sim \mathcal{U}([\text{SNR}_{\min}, \text{SNR}_{\max}])} \left[\left\| \mathbf{z}_0 - \hat{\mathbf{z}}_\theta(\alpha_{t(v)} \mathbf{z}_0 + \sigma_{t(v)} \boldsymbol{\xi}, t(v)) \right\|^2 \right] \end{aligned}$$

where $\hat{\mathbf{z}}_\theta$ is a neural network that predicts the denoised latent code \mathbf{z}_0 from the noised latent code \mathbf{z}_t , and $t(v)$ is the function that computes the time t from the SNR v .

- So, the overall loss that we want to minimize is:

$$E_{\mathbf{x} \sim p_{\text{data}}, \mathbf{z}_0 \sim q_\phi(\mathbf{z}_0|\mathbf{x})}$$

References

- [Khu20] Pramook Khungurn. Variational autoencoder. <https://pkhungurn.github.io/notes/notes/ml/vae/index.html>, 2020. Accessed: 2023-06-13.
- [Khu22] Pramook Khungurn. Variational diffusion models. <https://pkhungurn.github.io/notes/notes/ml/variational-diffusion-models/variational-diffusion-models.pdf>, 2022. Accessed: 2023-06-13.
- [KSPH21] Diederik P. Kingma, Tim Salimans, Ben Poole, and Jonathan Ho. Variational diffusion models, 2021.
- [SSDK⁺21] Yang Song, Jascha Sohl-Dickstein, Diederik P Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. Score-based generative modeling through stochastic differential equations. In *International Conference on Learning Representations*, 2021.
- [VK21] Arash Vahdat and Jan Kautz. Nvae: A deep hierarchical variational autoencoder, 2021.
- [VKK21] Arash Vahdat, Karsten Kreis, and Jan Kautz. Score-based generative modeling in latent space, 2021.