

Progressive Distillation of Diffusion Models

Pramook Khungurn

January 25, 2023

This note is written as I read “Progressive Distillation for Fast Sampling of Diffusion Models” by Salimans and Ho [SH22].

1 Introduction

- DDPMs can generate high quality samples. However, they are slow at sampling. Vanilla DDPMs (circa 2021) takes hundreds and thousands of models evaluations to sample one data item.
 - Note, however, that improvements by Karras et al. [KAAL22], for example, reduces the number of sampling steps to tens (20 to 50) in practice.
- The paper mentions that sampling a DDPM can be fast in “strongly conditioned settings.”
 - Text-to-speech.
 - Image super-resolution.
 - Classifier-guided sampling [DN21].
- It is the slowest when less conditioning information is available.
 - Unconditional sampling.
 - Class conditional sampling.
- The paper claims to make three contributions.
 - A new parameterizations of DDPMs that make them more stable when sampling with a few steps.
 - A method to distill a deterministic sampler that use many steps to one that uses a few steps.
 - Showing that the distillation process does not take more time than the process to train the original model.
- The core algorithm’s specification.
 - **Input.** A pre-trained DDPM that takes N steps to sample.
 - **Output.** A new DDPM that takes $N/2$ steps to sample.
- Progressive distillation = keep applying the core algorithm until it takes a few, say 4, steps to sample.
- Summary of results.
 - Start out with SOTA DDPM that takes 8192 steps to sample.
 - Distill it down to 4 steps.
 - Perceptual quality does not suffer much: FID of 3.0 on CIFAR-10 in 4 steps.

2 Background

- The paper uses the notations introduced in Kingma et al.’s “Variational Diffusion Models” paper [KSPH21].
- A data item is denoted by \mathbf{x} . The data distribution is denoted by p_{data} .
- Diffusion models work on **latent variables** $\{\mathbf{z}_t : t \in [0, 1]\}$. The latent variables form a Gaussian process that evolves over time according to the **forward process** $q(\mathbf{z}_t|\mathbf{x})$ where

$$q(\mathbf{z}_t|\mathbf{x}) = \mathcal{N}(\mathbf{z}_t; \alpha_t \mathbf{x}, \sigma_t^2 I)$$

where α_t and σ_t are differentiable functions of t , collectively known as the **noise schedule**.

- Let $\lambda_t = \log(\alpha_t^2/\sigma_t^2)$. We call it the **signal-to-noise ratio (SNR)**. We require that the SNR decreases monotonically with t and that

$$q(\mathbf{z}_t|\mathbf{z}_s) = \mathcal{N}\left(\mathbf{z}_t; \frac{\alpha_t}{\alpha_s} \mathbf{z}_s; \sigma_{t|s}^2 I\right)$$

where $0 \leq s < t \leq 1$ and

$$\sigma_{t|s}^2 = (1 - e^{\lambda_t - \lambda_s}) \sigma_t^2.$$

- We train a neural network $\hat{\mathbf{x}}_{\theta}$ so that $\hat{\mathbf{x}}_{\theta}(\mathbf{z}_t, \lambda_t) \approx \mathbf{x}$. In other words, $\hat{\mathbf{x}}_{\theta}(\mathbf{z}_t, \lambda_t)$ denoise \mathbf{z}_t back to \mathbf{x} . The network is trained with the following loss:

$$E_{\mathbf{x} \sim p_{\text{data}}, t \sim \mathcal{U}([0, 1]), \mathbf{z}_t \sim \mathcal{N}(\alpha_t \mathbf{x}, \sigma_t^2 I)} [w(\lambda_t) \|\hat{\mathbf{x}}_{\theta}(\mathbf{z}_t, \lambda_t) - \mathbf{x}\|^2].$$

Here, $\mathcal{U}([0, 1])$ is the uniform distribution over the interval $[0, 1]$. The weighting function $w(\lambda_t)$ will be discussed later.

- We can deduce that, for any $0 \leq s < t \leq 1$,

$$q(\mathbf{z}_s|\mathbf{z}_t, \mathbf{x}) = \mathcal{N}(\mathbf{z}_s; \tilde{\boldsymbol{\mu}}_{s|t}(\mathbf{z}_t, \mathbf{x}), \tilde{\sigma}_{s|t}^2 I)$$

where

$$\begin{aligned} \tilde{\boldsymbol{\mu}}_{s|t}(\mathbf{z}_t, \mathbf{x}) &= e^{\lambda_t - \lambda_s} \frac{\alpha_s}{\alpha_t} \mathbf{z}_t + (1 - e^{\lambda_t - \lambda_s}) \alpha_s \mathbf{x}, \\ \tilde{\sigma}_{s|t}^2 &= (1 - e^{\lambda_t - \lambda_s}) \sigma_s^2. \end{aligned}$$

- We can use the above equation as a basis for the **ancestral sampling algorithm**.
 1. We have a sequence of times $0 = t_0 < t_1 < t_2 < \dots < t_K = 1$.
 2. We start with $\mathbf{z}_1 \sim \mathcal{N}(0, I)$.
 3. Given that we have a sample at time t and we want sample at another time $s < t$, then,

$$\mathbf{z}_s = \tilde{\boldsymbol{\mu}}_{s|t}(\mathbf{z}_t, \hat{\mathbf{x}}_{\theta}(\mathbf{z}_t, \lambda_t)) + \sqrt{(\tilde{\sigma}_{s|t}^2)^{1-\gamma} (\sigma_{t|s}^2)^{\gamma}} \boldsymbol{\xi}$$

where $\boldsymbol{\xi} \sim \mathcal{N}(\mathbf{0}, I)$, and γ is a hyperparameter that controls how much noise is added during sampling [ND21].

- Alternatively, \mathbf{z}_t satisfies the following **probability flow ODE**:

$$d\mathbf{z}_t = \left(f(\mathbf{z}_t, t) - \frac{1}{2}g(t)^2 \nabla_{\mathbf{z}} \log \hat{p}_{\boldsymbol{\theta}}(\mathbf{z}_t) \right) dt$$

where

$$\begin{aligned} \nabla_{\mathbf{z}} \log \hat{p}_{\boldsymbol{\theta}}(\mathbf{z}_t) &= \frac{\alpha_t \hat{\mathbf{x}}_{\boldsymbol{\theta}}(\mathbf{z}_t, \lambda_t) - \mathbf{z}_t}{\sigma_t^2}, \\ f(\mathbf{z}_t, t) &= \frac{d \log \alpha_t}{dt} \mathbf{z}_t, \\ g(t)^2 &= \frac{d\sigma_t^2}{dt} - 2 \frac{d \log \alpha_t}{dt} \sigma_t^2. \end{aligned}$$

Sampling can then be done by simulating the ODE backward in time from $t = 1$ to $t = 0$.

- The DDIM sampler proposed by Song et al. [SME20] is given by

$$\begin{aligned} \mathbf{z}_s &= \alpha_s \hat{\mathbf{x}}_{\boldsymbol{\theta}}(\mathbf{z}_t) + \sigma_s \frac{\mathbf{z}_t - \alpha_t \hat{\mathbf{x}}_{\boldsymbol{\theta}}(\mathbf{z}_t)}{\sigma_t} \\ &= e^{(\lambda_t - \lambda_s)/2} (\alpha_s / \alpha_t) \mathbf{z}_t + (1 - e^{(\lambda_t - \lambda_s)/2}) \alpha_s \hat{\mathbf{x}}_{\boldsymbol{\theta}}(\mathbf{z}_t). \end{aligned}$$

The paper claims that this can be understood in terms of integrating the probability flow ODE.

3 Progressive Distillation

- The core algorithm receives as input a **teacher model** that is trained in the standard way.

STANDARD-DIFFUSION-TRAINING

```

1  while not converged
2       $\mathbf{x} \sim p_{\text{data}}$ 
3       $t \sim \mathcal{U}([0, 1])$ 
4       $\boldsymbol{\xi} \sim \mathcal{N}(\mathbf{0}, I)$ 
5       $\mathbf{z}_t \leftarrow \alpha_t \mathbf{x} + \sigma_t \boldsymbol{\theta}$ 
6       $\lambda_t \leftarrow \log(\alpha_t^2 / \sigma_t^2)$ 
7       $L_{\boldsymbol{\theta}} \leftarrow w(\lambda_t) \|\mathbf{x} - \hat{\mathbf{x}}_{\boldsymbol{\theta}}(\mathbf{z}_t, \lambda_t)\|_2^2$ 
8      Update  $\boldsymbol{\theta}$  according to  $\nabla_{\boldsymbol{\theta}} L_{\boldsymbol{\theta}}$ .
```

- The result of the distillation process is a **student model**, which has the same architecture as the teacher model.
- The distillation process proceeds in iterations.
 - Each iteration halves the number of sampling steps required.
 - At the start of each iteration, the student model is initialized with the parameters of the teacher model.
 - At the end of each iteration, the student model becomes the teacher model of the next round.
 - The training process in each iteration is similar to the standard training. However, the main difference is the target that $\hat{\mathbf{x}}_{\boldsymbol{\theta}}$ is asked to approximate from \mathbf{z}_t .
 - * For the standard training, the target is \mathbf{x} .

- * For the distillation process, the target is the denoised data that would have been predicted by the teacher in two distillation steps.
- Another difference is that the distillation process always work in discrete time instead of continuous time like the standard training.
- Here’s the distillation algorithm.

PROGRESSIVE-DISTILLATION(η, N)

```

    //  $\eta$  denotes the parameters of the trained teacher model.
    //  $N$  is the number of iterations that the teacher model takes to sample.
1  for  $K$  iterations
2       $N \leftarrow N/2$ 
3       $\theta \leftarrow \eta$ 
4      while not converged
5           $\mathbf{x} \sim p_{\text{data}}$ 
6           $i \sim \mathcal{U}(\{1, 2, \dots, N\})$ 
7           $t \leftarrow i/N$ 
8           $\xi \sim \mathcal{N}(\mathbf{0}, I)$ 
          // 2 steps of DDIM sampling with teacher model
9           $t' \leftarrow t - 0.5/N$ ;    $t'' \leftarrow t - 1/N$ 
10          $\lambda_t \leftarrow \log(\alpha_t^2/\sigma_t^2)$ ;    $\lambda_{t'} \leftarrow \log(\alpha_{t'}^2/\sigma_{t'}^2)$ 
11          $\mathbf{z}_{t'} \leftarrow \alpha_{t'} \hat{\mathbf{x}}_{\eta}(\mathbf{z}_t, \lambda_t) + \frac{\sigma_{t'}}{\sigma_t} (\mathbf{z}_t - \alpha_t \hat{\mathbf{x}}_{\eta}(\mathbf{z}_t, \lambda_t))$ 
12          $\mathbf{z}_{t''} \leftarrow \alpha_{t''} \hat{\mathbf{x}}_{\eta}(\mathbf{z}_{t'}, \lambda_{t'}) + \frac{\sigma_{t''}}{\sigma_{t'}} (\mathbf{z}_{t'} - \alpha_{t'} \hat{\mathbf{x}}_{\eta}(\mathbf{z}_{t'}, \lambda_{t'}))$ 
13          $\tilde{\mathbf{x}} \leftarrow \frac{\mathbf{z}_{t''} - (\sigma_{t''}/\sigma_t)\mathbf{z}_t}{\alpha_{t''} - (\sigma_{t''}/\sigma_t)\alpha_t}$ 
14          $L_{\theta} \leftarrow w(\lambda_t) \|\tilde{\mathbf{x}} - \hat{\mathbf{x}}_{\theta}(\mathbf{z}_t, \lambda_t)\|^2$ 
15         Update  $\theta$  according to  $\nabla_{\theta} L_{\theta}$ 
          // Make the converged student the teacher of the next round.
16      $\eta \leftarrow \theta$ 

```

- The paper notes that using the value predicted by the teacher in two steps as a target makes the prediction task for the student model much easier.
 - If using \mathbf{x} as the target, the model must predict the average of possible \mathbf{x} value, which produces blurry predictions.
 - However, the value predicted by the teacher model is completely determined by the teacher model and \mathbf{z}_t . So, the prediction is sharp.

As a result, the student model can make progress much faster than the vanilla teacher model can during sampling.

4 Parameterization and Training Loss

- This section is about the details of the models: α_t , σ_t , $w(\lambda_t)$, and what $\hat{\mathbf{x}}_{\theta}$ is.
- The paper assumes a variance preserving model: $\sigma_t^2 = 1 - \alpha_t^2$.
- Cosine schedule is used: $\alpha_t = \cos(0.5\pi t)$.

4.1 Model Parameterization

- Many works on DDPM parameterize $\hat{\mathbf{x}}_{\theta}$ by requiring it to predict the noise ξ that is used to make \mathbf{z}_t from \mathbf{x} . In other words, our network is $\hat{\xi}_{\theta}$, and we take

$$\hat{\mathbf{x}}_{\theta}(\mathbf{z}_t, \lambda_t) = \frac{1}{\alpha_t}(\mathbf{z}_t - \sigma_t \hat{\xi}_{\theta}(\mathbf{z}_t, \lambda_t)).$$

- While the above parameterization might work well for training the original model, the paper argues that this does not work well with distillation.
 - As distillation progresses, we increasingly evaluate at lower and lower signal-to-noise ratios.
 - As $\alpha_t \rightarrow 0$, the changes made by $\hat{\xi}_{\theta}$ gets amplified drastically. This is because

$$\hat{\mathbf{x}}_{\theta}(\mathbf{z}_t, \lambda_t) = \frac{1}{\alpha_t}(\mathbf{z}_t - \sigma_t \hat{\xi}_{\theta}(\mathbf{z}_t, \lambda_t)).$$

So, changes are multiplied by $1/\alpha_t$.

- If we distill down to a single sampling step, the input to the model is pure noise ξ , and the distilled $\hat{\xi}_{\theta}$ would have to predict ξ . The link between ξ -prediction and \mathbf{x} -prediction breaks down completely.
- So, for distillation to work, we need to parameterize the model so that the prediction $\hat{\mathbf{x}}_{\theta}$ remains stable as the SNR varies. The paper tried a number of combinations and found all of them to work well.
 - Predicting \mathbf{x} directly.
 - Predicting both \mathbf{x} and ξ , resulting in $\tilde{\mathbf{x}}$ and $\tilde{\xi}$, and then merging via

$$\hat{\mathbf{x}} = \sigma_t^2 \tilde{\mathbf{x}} + \alpha_t(\mathbf{z}_t - \sigma_t \tilde{\xi}).$$

Note that $\sigma_t^2 \tilde{\mathbf{x}} \approx \sigma_t^2 \mathbf{x}$, and

$$\alpha_t(\mathbf{z}_t - \sigma_t \tilde{\xi}) = \alpha_t(\alpha_t \mathbf{x} + \sigma_t \theta - \sigma_t \tilde{\xi}) \approx \alpha_t^2 \mathbf{x}.$$

Because $\alpha_t^2 + \sigma_t^2 = 1$, we have that $\sigma_t^2 \tilde{\mathbf{x}} + \alpha_t(\mathbf{z}_t - \sigma_t \tilde{\xi}) \approx \sigma_t^2 \mathbf{x} + \alpha_t^2 \mathbf{x}$ is an interpolation between the \mathbf{x} -prediction and the implied \mathbf{x} from ξ -prediction.

- Predicting $\mathbf{v} := \alpha_t \xi - \sigma_t \mathbf{x}$ and setting $\hat{\mathbf{x}} := \alpha_t \mathbf{z}_t - \sigma_t \hat{\mathbf{v}}_{\theta}(\mathbf{z}_t, \lambda_t)$.
- The paper also tried all the above combinations when training the original models, and it found that they worked well too.

4.2 Weighting Function

- In the 2020 work by Ho et al. [HJA20], they sample the time uniformly and compute the squared L2-distance between ξ and $\hat{\xi}_{\theta}(\mathbf{z}_t, \lambda_t)$. One can deduce that

$$L_{\theta} = \|\xi - \hat{\xi}_{\theta}(\mathbf{z}_t, \lambda_t)\|^2 = \frac{\alpha_t^2}{\sigma_t^2} \|\mathbf{x} - \hat{\mathbf{x}}_{\theta}(\mathbf{z}_t, \lambda_t)\|^2.$$

This is called the **SNR** weighting scheme. This means that the loss gives zero weight to data with zero SNR. This is not a suitable loss for distillation.

- The paper tried the following choices of weighting scheme.

- **Truncated SNR:** $L_{\theta} = \max(\|\mathbf{x} - \hat{\mathbf{x}}\|^2, \|\boldsymbol{\xi} - \hat{\boldsymbol{\xi}}\|^2) = \max(\frac{\alpha_t^2}{\sigma_t^2}, 1)\|\mathbf{x} - \hat{\mathbf{x}}\|^2$.
- **SNR+1:** $L_{\theta} = \|\mathbf{v} - \hat{\mathbf{v}}\| = (1 + \frac{\alpha_t^2}{\sigma_t^2})\|\mathbf{x} - \hat{\mathbf{x}}\|^2$.

It found that both worked well.

- The authors did not put much attention on how the time is sampled. They just sampled it uniformly randomly from $[0, 1]$.

5 Experiments

5.1 Model Parameterization and Training Loss

- The paper tried the 4 combinations of model parameterization and 3 combination of weighting schemes.
 - One case lead to training divergence: predicting $\boldsymbol{\xi}$ and using the truncated SNR weighting scheme.
 - In all other cases, the numbers are similar.
- The paper observed that predicting \mathbf{v} is the most stable method because it makes the DDIM step-sizes independent of the SNR. However, predicting \mathbf{x} gives the best empirical results.

5.2 Progressive Distillation

- The authors started with a teacher model trained on continous time. They then started distilling from $N = 8192$ (or $N = 1024$ for bigger models) down to 1.
- In each distillation iteration, they optimize the model for 50 000 parameter updates. The exceptions are when distilling to 2 or 1 step(s), which they took 100 000 updates.
- FID scores of distilled models increases very slowly until $N = 4$, after which they rises very rapidly. So, $N = 4$ seems to be a sweet spot between speed and quality.
- Undistilled models' performance degrades very fast after $N = 128$.

References

- [DN21] Prafulla Dhariwal and Alexander Nichol. Diffusion models beat gans on image synthesis. In M. Ranzato, A. Beygelzimer, Y. Dauphin, P.S. Liang, and J. Wortman Vaughan, editors, *Advances in Neural Information Processing Systems*, volume 34, pages 8780–8794. Curran Associates, Inc., 2021.
- [HJA20] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *CoRR*, abs/2006.11239, 2020.
- [KAAL22] Tero Karras, Miika Aittala, Timo Aila, and Samuli Laine. Elucidating the design space of diffusion-based generative models, 2022.
- [KSPH21] Diederik P. Kingma, Tim Salimans, Ben Poole, and Jonathan Ho. Variational diffusion models, 2021.
- [ND21] Alex Nichol and Prafulla Dhariwal. Improved denoising diffusion probabilistic models, 2021.
- [SH22] Tim Salimans and Jonathan Ho. Progressive distillation for fast sampling of diffusion models. *CoRR*, abs/2202.00512, 2022.
- [SME20] Jiaming Song, Chenlin Meng, and Stefano Ermon. Denoising diffusion implicit models, 2020.