

# Diffusion Models and Inverse Problems

Pramook Khungurn

March 19, 2025

This note is written as I read the survey paper “A Survey on Diffusion Models for Inverse Problems” by Dara et al. [DCL<sup>+</sup>24].

## 1 Problem Setting

- A data item is denoted by  $\mathbf{x} \in \mathbb{R}^n$ .
- We are interested in a random variable  $\mathbf{X}$  whose values are data items. The probability distribution of  $\mathbf{X}$  is denoted by  $p_{\mathbf{X}}$ .
- Given a data item  $\mathbf{x}$ , we extract from it a **measurement**  $\mathbf{y} \in \mathbb{R}^m$ .
- We model the measurement process with a **measurement model**, which is also called a **forward model** or a **corruption model** in literature. That is, we establish a random variable  $\mathbf{Y}$  whose values are measurements according to the following equation.

$$\mathbf{Y} = \mathcal{A}(\mathbf{X}) + \sigma_{\mathbf{y}}\mathbf{Z}$$

where  $\mathcal{A} : \mathbb{R}^n \rightarrow \mathbb{R}^m$  is a deterministic function and  $\mathbf{Z} \sim \mathcal{N}(0, I)$ . The function  $\mathcal{A}$  is often called the **forward measurement operator** or just the **measurement operator**.

- In an inverse problem, we are given a measurement  $\mathbf{y}$ . The goal is to produce a data item  $\mathbf{x}$  that could have produced  $\mathbf{y}$  according to the measurement model.
- Examples of inverse problems. Here,  $\sigma_{\mathbf{y}} = 0$  unless noted otherwise.
  - **Denoising.**  $\mathcal{A}$  is the identity function, and  $\sigma_{\mathbf{y}} > 0$ . This results in the measurement  $\mathbf{y}$  being  $\mathbf{x}$  with some added noise.
  - **Inpainting.**  $\mathcal{A}$  zeros out some components  $\mathbf{x}$ .
  - **Compressed sensing.**  $\mathcal{A}(\mathbf{x}) = A\mathbf{x}$  where  $A$  is a matrix with entries sampled from a Gaussian random variable.
  - **Signal recovery from convolution.**  $\mathcal{A}(\mathbf{x}) = \mathbf{x} * \mathbf{k}$  where  $*$  denotes convolution. **Deblurring** and **super-resolution** are examples of signal recovery from convolution.

Note that all of the examples above are *linear* in the sense that  $\mathcal{A}(\mathbf{x}) = A\mathbf{x}$  for some matrix  $A \in \mathbb{R}^{m \times n}$ . Moreover,  $m = n$  in all of them. However, these constraints need not hold in general. Here is an example of a nonlinear inverse problem.

- **Phase retrieval.**  $\mathcal{A}(\mathbf{x}) = |\mathcal{F}(\mathbf{x})|$  where  $\mathcal{F}$  is the discrete Fourier transform operator, and  $|\cdot|$  denotes the norm of a complex number. Here, the measurement is the norms of the Fourier coefficients, which means that the phase information has been thrown away.

- **Decompression.**  $\mathcal{A}$  is a non-linear function that compresses  $\mathbf{x}$ . For example,  $\mathbf{x}$  can be an image, and  $\mathcal{A}$  is the JPEG compression algorithm.
- One of the characteristic of the above problems are that perfect recovery is impossible.
  - In other words, these problems are *ill-posed*.
- As result, we almost always have say exactly what type of recovery we want.
  - **Posterior sampling.** We do not care about the exact  $\mathbf{x}$ , but we just want to sample from the **posterior** distribution  $p_{\mathbf{X}}(\mathbf{x}|\mathbf{Y} = \mathbf{y})$ .
    - \* For brevity, we shall write this as just  $p(\mathbf{x}|\mathbf{y})$ .
  - **Maximum a posterior (MAP) estimation.** We look for  $\mathbf{x}$  that maximizes  $p(\mathbf{x}|\mathbf{y})$ .
  - **Minimum mean square error (MMSE) estimation.** We seek  $\mathbf{x}^*$  that minimizes

$$\mathcal{L}(\mathbf{x}^*) = E_{\mathbf{x} \sim p(\mathbf{x}|\mathbf{y})}[\|\mathbf{x} - \mathbf{x}^*\|^2].$$

This results in the expectation  $E[\mathbf{X}|\mathbf{Y} = \mathbf{y}]$  (or just simply  $E[\mathbf{x}|\mathbf{y}]$ ).

- There are limitations to MAP and MMSE estimations [BM18, DM24].
  - They can suffer from “regression to the mean.”
  - This means that the predicted  $\mathbf{x}$  may lack important details and be outside the desire solution space.

So, in this note, the focus is on posterior sampling.

- As the title of this note implies, we are interested in using diffusion models to sample from  $p(\mathbf{x}|\mathbf{y})$ .
- A direct way to do so is to just train a conditional diffusion model to do this job.
- Indeed, there are many conditional diffusion models for specific inverse problems such as:
  - super-resolution [LYC<sup>+</sup>21, SHC<sup>+</sup>21],
  - deblurring [WDT<sup>+</sup>21],
  - inpainting [SCC<sup>+</sup>22], and
  - image restoration [SCC<sup>+</sup>22, LGZ<sup>+</sup>23a, LGZ<sup>+</sup>23b].

Even ControlNet [ZRA23] and similar approches can be thought of an instance of solving an inverse problem.

- However, these approches require training a new diffusion model from scratch or fine-tuning existing ones. This can be quite expensive.
- Instead, we focus on the **unsupervised approach**, in which we take an existing diffusion model and use it to sample from  $p(\mathbf{x}|\mathbf{y})$  without any training or fine-tuning.
- Let us first see how using a diffusion model can be useful. To see this we need to mathematically examine posterior sampling. According to Bayes’ rule, we have that

$$p(\mathbf{x}|\mathbf{y}) = p(\mathbf{X} = \mathbf{x}|\mathbf{Y} = \mathbf{y}) = \frac{p(\mathbf{Y} = \mathbf{y}|\mathbf{X} = \mathbf{x})p(\mathbf{X} = \mathbf{x})}{p(\mathbf{Y} = \mathbf{y})} = \frac{p_{\mathbf{Y}}(\mathbf{y}|\mathbf{x})p_{\mathbf{X}}(\mathbf{x})}{p_{\mathbf{Y}}(\mathbf{y})}.$$

We generally do not care about  $p_{\mathbf{Y}}(\mathbf{y})$  because  $\mathbf{y}$  is already given to us. So, we usually just write

$$p(\mathbf{x}|\mathbf{y}) \propto p(\mathbf{y}|\mathbf{x})p_{\mathbf{X}}(\mathbf{x}).$$

Our task thus becomes sampling from the unnormalized distribution  $p(\mathbf{y}|\mathbf{x})p_{\mathbf{X}}(\mathbf{x})$ . The distribution  $p_{\mathbf{X}}(\mathbf{x})$  is called the **prior** and  $p(\mathbf{y}|\mathbf{x})$  is called the **likelihood**.

- A pre-trained diffusion model can be useful because we can use it as a prior  $p_{\mathbf{X}}$ . However, there are several difficulty in this direction.
- First, given a sample  $\mathbf{x}$ , a pre-trained diffusion model does not allow us evaluate  $p_{\mathbf{X}}(\mathbf{x})$  directly.
- Second, we need to find a way to somehow approximate  $p(\mathbf{y}|\mathbf{x})$  or something related to it with the diffusion model.

## 2 Background

### 2.1 Diffusion Processes

- Given a data distribution  $p_{\mathbf{X}}$ , we can construct a stochastic process  $\{\mathbf{X}_t : 0 \leq t \leq T\}$  such that  $\mathbf{X}_0 \sim p_{\mathbf{X}}$  and  $\mathbf{X}_T$  is approximately distributed according to  $\mathcal{N}(\mathbf{0}, \sigma_T^2 I)$  for some  $\sigma_T > 0$ . This is done through a stochastic differential equation (SDE):

$$d\mathbf{X}_t = \mathbf{f}(\mathbf{X}_t, t) dt + g(t) d\mathbf{W}_t \quad (1)$$

where  $W_t$  is the standard Brownian motion in  $\mathbb{R}^n$ . The stochastic process we desire is the solution to the above SDE with the initial condition  $\mathbf{X}_0 \sim p_{\mathbf{X}}$ .

- Let us give names to functions inside the above SDE.
  - The function  $\mathbf{f} : \mathbb{R}^n \times \mathbb{R} \rightarrow \mathbb{R}^n$  is called the **drift coefficient**.
  - The function  $g : \mathbb{R} \rightarrow \mathbb{R}$  is called the **diffusion coefficient**.
- We will be working with the distribution of  $\mathbf{X}_t$ , which is denoted by  $p_{\mathbf{X}_t}$  under standard notation. However, for convenience and brevity, let us use  $p_t$  instead of  $p_{\mathbf{X}_t}$ . In other words,

$$p_t(\mathbf{x}_t) = p_{\mathbf{X}_t}(\mathbf{x}_t) = p(\mathbf{X}_t = \mathbf{x}_t).$$

- According to Anderson [And82], we can sample  $p_0$  by solving the reverse SDE

$$d\mathbf{X}_t = \left( \mathbf{f}(\mathbf{X}_t, t) - g^2(t) \nabla_{\mathbf{X}_t} \log p_t(\mathbf{X}_t) \right) dt + g(t) d\overline{\mathbf{W}}_t \quad (2)$$

from  $t = T$  to  $t = 0$ , initializing  $\mathbf{X}_T$  by sampling from  $p_T$ , an isotropic Gaussian distribution. Here,  $\overline{\mathbf{W}}_t$  is the standard Brownian motion that runs backwards in time.

- Song et al. observes that the following ODE,

$$\frac{d\mathbf{x}_t}{dt} = \mathbf{f}(\mathbf{x}_t, t) - \frac{g^2(t)}{2} \nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t), \quad (3)$$

satisfies the same Fokker–Planck equation as the SDE (1) [SSDK<sup>+</sup>21]. Hence, we also have a deterministic sampling scheme to sample from  $p_{\mathbf{X}}$ .

- There are two variants of the SDE in Equation 1 that are widely used.
  - **Variance exploding SDE.**
    - \* We set  $\mathbf{f}(\mathbf{X}_t, t) = \mathbf{0}$ .
    - \* The SDE (1) then boils down to the following relationship between random variables.

$$\mathbf{X}_t = \mathbf{X}_0 + \sigma_t \mathbf{Z}$$

where  $\mathbf{X}_0 \sim p_0 = p_{\mathbf{X}}$ ,  $\mathbf{Z} \sim \mathcal{N}(\mathbf{0}, I)$ , and

$$\sigma_t = \sqrt{\int_0^t g^2(u) du}.$$

- \* A typical noise schedule is  $\sigma_t = t$ , which corresponds to  $g(t) = \sqrt{2t}$ , and  $\mathbf{X}_t = \mathbf{X}_0 + t\mathbf{Z}$ . This schedule is used in the famous EDM paper [KAAL22].

– **Variance preversing SDE.**

- \* We choose the drift and the diffusion coefficients so that the following relationship holds:

$$\mathbf{X}_t = \alpha_t \mathbf{X}_0 + \sigma_t \mathbf{Z}, \text{ and } \alpha_t^2 + \sigma_t^2 = 1.$$

- \* Again,  $\mathbf{X}_0 \sim p_0 = p_{\mathbf{X}}$  and  $\mathbf{Z} \sim \mathcal{N}(\mathbf{0}, I)$ .
- \* The functions  $\alpha_t : [0, T] \rightarrow \mathbb{R}$  and  $\sigma_t : [0, T] \rightarrow \mathbb{R}$  are both called “noise schedules.”
- \* An example of a variance preserving SDE is the following Ornstein–Uhlenbeck process,

$$d\mathbf{X}_t = -\mathbf{X}_t dt + \sqrt{2} d\mathbf{W}_t,$$

which gives  $\alpha_t = \exp(-t)$  and  $\sigma_t = \sqrt{1 - \exp(-2t)}$ .

- In general, we will be using an SDE of the form

$$d\mathbf{X}_t = f(t)\mathbf{X}_t dt + g(t) d\mathbf{W}_t$$

for some  $f : \mathbb{R} \rightarrow \mathbb{R}$ . This equation gives us the relation

$$\mathbf{X}_t = \alpha_t \mathbf{X}_0 + \sigma_t \mathbf{Z} \tag{4}$$

where  $\mathbf{X}_0 \sim p_0 = p_{\mathbf{X}}$  and  $\mathbf{Z} \sim \mathcal{N}(\mathbf{0}, I)$ . The noise schedules  $\alpha_t$  and  $\sigma_t$  are related to  $f(t)$  and  $g(t)$  as follows [KAAL22]:

$$\begin{aligned} \alpha_t &= \exp\left(\int_0^t f(u) du\right), \\ \sigma_t &= \sqrt{\int_0^t \frac{g^2(u)}{\alpha_u^2} du}, \\ f(t) &= \frac{\dot{\alpha}_t}{\alpha_t}, \\ g(t) &= \alpha_t \sqrt{2\sigma_t \dot{\sigma}_t}. \end{aligned}$$

Here,  $\dot{\alpha}_t$  and  $\dot{\sigma}_t$  are derivatives of  $\alpha_t$  and  $\sigma_t$  with respect to time.

## 2.2 Tweedie’s Formula and Diffusion Model Training

- **Theorem 1 (Tweedie’s formula.).** *Let  $\mathbf{X}$  and  $\mathbf{Y}$  be random variables such that  $\mathbf{Y} = \mathbf{X} + \sigma\mathbf{Z}$  where  $\mathbf{Z} \sim \mathcal{N}(\mathbf{0}, I)$ . It follows that*

$$E[\mathbf{X} | \mathbf{Y} = \mathbf{y}] = \mathbf{y} + \sigma^2 \nabla_{\mathbf{y}} \log p_{\mathbf{Y}}(\mathbf{y}).$$

- The expression  $\nabla_{\mathbf{y}} \log p_{\mathbf{Y}}(\mathbf{y})$  is called the **score** of  $p_{\mathbf{Y}}$ .
- Applying Tweedie’s formula to Equation (4), we have that

$$E[\alpha_t \mathbf{X}_0 | \mathbf{X}_t = \mathbf{x}_t] = \mathbf{x}_t + \sigma_t^2 \nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t).$$

In other words,

$$E[\mathbf{X}_0 | \mathbf{X}_t = \mathbf{x}_t] = \frac{\mathbf{x}_t + \sigma_t^2 \nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t)}{\alpha_t}.$$

Or, more concisely,

$$E[\mathbf{x}_0|\mathbf{x}_t] = \frac{\mathbf{x}_t + \sigma_t^2 \nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t)}{\alpha_t}. \quad (5)$$

This gives

$$\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t) = \frac{\alpha_t E[\mathbf{x}_0|\mathbf{x}_t] - \mathbf{x}_t}{\sigma_t^2}. \quad (6)$$

- A diffusion model is a neural network  $\mathbf{h}_\theta$  with parameters  $\theta$  such that  $\mathbf{h}_\theta(\mathbf{x}_t, t)$  approximates  $E[\mathbf{x}_0|\mathbf{x}_t]$ . It can be trained with the following **denoising score matching** objective

$$\mathcal{L}_{\text{DSM}}(\theta) = E_{t, \mathbf{x}_0 \sim p_{\mathbf{x}}, \mathbf{z} \sim \mathcal{N}(\mathbf{0}, I)} \left[ \left\| \mathbf{h}_\theta(\alpha_t \mathbf{x}_0 + \sigma_t \mathbf{z}, t) \right\|^2 \right].$$

- Let  $\hat{\mathbf{x}}_0(\mathbf{x}_t)$  denotes the expected  $\mathbf{x}_0$  given  $\mathbf{x}_t$  computed by the neural network.

$$\hat{\mathbf{x}}_0(\mathbf{x}_t) = \mathbf{h}_\theta(\mathbf{x}_t, t).$$

- It follows that we may estimate the score of  $p_t$  as follows:

$$\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t) \approx \frac{\alpha_t \hat{\mathbf{x}}_0(\mathbf{x}_t) - \mathbf{x}_t}{\sigma_t^2}. \quad (7)$$

### 2.3 Denoising Diffusion Implicit Model (DDIM)

- Given a stochastic process  $\{\mathbf{X}_t : 0 \leq t \leq T\}$  where Equation (4) holds, Song et al. gives an alternative process where  $\mathbf{X}_t$  can be generated [SME22].
  1. Sample  $\mathbf{X}_0 \sim p_0$ .
  2. Sample  $\mathbf{X}_T \sim p(\mathbf{X}_t|\mathbf{X}_0) = \mathcal{N}(\mathbf{X}_t; \alpha_T \mathbf{X}_0, \sigma_T^2 I)$ .
  3. Given that some  $t' > t$  has been sampled, sample  $\mathbf{X}_t$  according to  $p(\mathbf{X}_t|\mathbf{X}_0, \mathbf{X}_{t'})$ , which is given by

$$p(\mathbf{X}_t|\mathbf{X}_0, \mathbf{X}_{t'}) = \mathcal{N}\left(\mathbf{X}_t; \alpha_t \mathbf{X}_0 + \sqrt{\sigma_t^2 - \varsigma^2} \frac{\mathbf{X}_{t'} - \alpha_{t'} \mathbf{X}_0}{\sigma_{t'}}, \varsigma^2 I\right)$$

This process works for any value of  $\varsigma \geq 0$ .

- In particular, if  $\varsigma = 0$ , then we simply set

$$\mathbf{X}_t \leftarrow \alpha_t \mathbf{X}_0 + \sigma_t \frac{\mathbf{X}_{t'} - \alpha_{t'} \mathbf{X}_0}{\sigma_{t'}}. \quad (8)$$

In fact,  $\mathbf{X}_t$  for all  $0 < t < T$  is deterministic given we have sampled  $\mathbf{X}_0$  and  $\mathbf{X}_T$ .

- However, the process we just described cannot be used to sample from  $p_0$  because it requires being able to sample from  $p_0$  in the first place.
- Song et al. proposes a practical sampling process by changing Equation (8) to use  $E[\mathbf{X}_0|\mathbf{X}_{t'}]$  in place of  $\mathbf{X}_0$ .

$$\mathbf{X}_t \leftarrow E[\mathbf{X}_0|\mathbf{X}_{t'}] + \sigma_t \frac{\mathbf{X}_{t'} - \alpha_{t'} E[\mathbf{X}_0|\mathbf{X}_{t'}]}{\sigma_{t'}} \quad (9)$$

Now, there is no dependency on  $\mathbf{X}_0$  any more. In fact, we can use a diffusion model  $\mathbf{h}_\theta$  to estimate  $E[\mathbf{X}_0|\mathbf{X}_{t'}]$ .

- The RHS of Equation (9) is an important operation, so let's give it a name. Let

$$\text{UncondDDIM}(\mathbf{x}_{t'}, \hat{\mathbf{x}}_0, t', t) = \alpha_t \hat{\mathbf{x}}_0 + \sigma_t \frac{\mathbf{x}_{t'} - \alpha_{t'} \hat{\mathbf{x}}_0}{\sigma_{t'}} = \left( \alpha_t - \frac{\alpha_{t'}}{\sigma_{t'}} \right) \hat{\mathbf{x}}_0 + \frac{\sigma_t}{\sigma_{t'}} \mathbf{x}_{t'}$$

- The **DDIM sampling** algorithm uses the UncondDDIM operation to sample from  $p_0$ . It requires a series of times  $0 = t_0 < t_1 < t_2 < \dots < t_K = T$ , and it goes as follows.
  - 1: Sample  $\mathbf{x}_K \sim \mathcal{N}(\mathbf{0}, \sigma_T^2 I)$ .
  - 2: **for**  $k \leftarrow K$  **downto** 1 **do**
  - 3:    $\hat{\mathbf{x}}_0 \leftarrow \mathbf{h}_\theta(\mathbf{x}_k, t_k)$
  - 4:    $\mathbf{x}_{k-1} \leftarrow \text{UncondDDIM}(\mathbf{x}_k, \hat{\mathbf{x}}_0, t_k, t_{k-1})$
  - 5: **end for**
  - 6: **return**  $\mathbf{x}_0$

## 2.4 Conditional Sampling

- The goal of inverse problem is to sample from  $p_0(\cdot|\mathbf{y})$  assuming  $\mathbf{Y} = \mathcal{A}(\mathbf{X}_0) + \sigma_{\mathbf{y}} \mathbf{Z}$  with  $\mathbf{Z} \sim \mathcal{N}(\mathbf{0}, I)$ .
- We can use the framework for diffusion models to do the job. In particular, we can sample from  $p_0(\cdot|\mathbf{y})$  by solving a variant of Equation 2:

$$d\mathbf{X}_t = \left( \mathbf{f}(\mathbf{X}_t, t) - g^2(t) \nabla_{\mathbf{X}_t} \log p_t(\mathbf{X}_t | \mathbf{Y} = \mathbf{y}) \right) dt + g(t) d\bar{\mathbf{W}}_t.$$

Equivalently, we can also solve a similar variant of the probability flow ODE (3):

$$\frac{d\mathbf{x}_t}{dt} = \mathbf{f}(\mathbf{x}_t, t) - \frac{g^2(t)}{2} \nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t | \mathbf{Y} = \mathbf{y}). \quad (10)$$

- One can see that we require the **conditional score**  $\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t | \mathbf{Y} = \mathbf{y})$ .
  - For brevity, we shall write it as  $\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t | \mathbf{y})$ .
- For supervised approach, we train  $\mathbf{h}_\theta(\mathbf{x}_t, t, \mathbf{y})$  so that it approximates  $E[\mathbf{X}_0 | \mathbf{X}_t = \mathbf{x}_t, \mathbf{Y} = \mathbf{y}]$  with the following **conditional denoising score matching** (CDSM) loss

$$\mathcal{L}_{\text{CDSM}}(\theta) = E_{t, \mathbf{x} \sim p_{\mathbf{x}}, \mathbf{z}_{\mathbf{x}} \sim \mathcal{N}(\mathbf{0}, I_n), \mathbf{z}_{\mathbf{y}} \sim \mathcal{N}(\mathbf{0}, I_m)} \left[ \left\| \mathbf{h}_\theta(\alpha_t \mathbf{x} + \sigma_t \mathbf{z}_{\mathbf{x}}, t, \mathcal{A}(\mathbf{x}) + \sigma_{\mathbf{y}} \mathbf{z}_{\mathbf{y}}) - \mathbf{x} \right\|^2 \right].$$

Then, we can approximate  $E[\mathbf{X}_0 | \mathbf{X}_t = \mathbf{x}_t, \mathbf{Y} = \mathbf{y}]$  with

$$\hat{\mathbf{x}}(\mathbf{x}_t | \mathbf{y}) = \mathbf{h}_\theta(\mathbf{x}_t, t, \mathbf{y})$$

and  $\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t | \mathbf{y})$  with

$$\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t | \mathbf{y}) \approx \frac{\alpha_t \hat{\mathbf{x}}_0(\mathbf{x}_t | \mathbf{y}) - \mathbf{x}_t}{\sigma_t^2}.$$

- However, in unsupervised approaches, we only have access to a vanilla diffusion model  $\mathbf{h}_\theta(\mathbf{x}_t, t)$ , not a conditional diffusion model  $\mathbf{h}_\theta(\mathbf{x}_t, t, \mathbf{y})$ . So, we have to do something else.
- Because

$$p(\mathbf{X}_t = \mathbf{x}_t | \mathbf{Y} = \mathbf{y}) = \frac{p(\mathbf{Y} = \mathbf{y} | \mathbf{X}_t = \mathbf{x}_t) p(\mathbf{X}_t = \mathbf{x}_t)}{p(\mathbf{Y} = \mathbf{y})}.$$

It follows that

$$\nabla_{\mathbf{x}_t} \log p(\mathbf{X}_t = \mathbf{x}_t | \mathbf{Y} = \mathbf{y}) = \nabla_{\mathbf{x}_t} \log p(\mathbf{Y} = \mathbf{y} | \mathbf{X}_t = \mathbf{x}_t) + \nabla_{\mathbf{x}_t} \log p(\mathbf{X}_t = \mathbf{x}_t) - \nabla_{\mathbf{x}_t} \log p(\mathbf{Y} = \mathbf{y})$$

In other words,

$$\nabla_{\mathbf{x}_t} \log p_{\text{ask}}(\mathbf{x}_t | \mathbf{y}) = \nabla_{\mathbf{x}_t} \log p(\mathbf{y} | \mathbf{x}_t) + \nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t) \quad (11)$$

- This means that we can rewrite the reverse time SDE as

$$d\mathbf{X}_t = \left( \mathbf{f}(\mathbf{X}_t, t) - g^2(t) (\nabla_{\mathbf{x}_t} \log p(\mathbf{y}|\mathbf{X}_t) + \nabla_{\mathbf{x}_t} \log p_t(\mathbf{X}_t)) \right) dt + g(t) d\bar{\mathbf{W}}_t$$

and the probability flow ODE as

$$\frac{d\mathbf{x}_t}{dt} = \mathbf{f}(\mathbf{x}_t, t) - \frac{g^2(t)}{2} (\nabla_{\mathbf{x}_t} \log p(\mathbf{y}|\mathbf{x}_t) + \nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t)).$$

- We already know how to approximate  $\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t)$  with a diffusion model. So, all we have left to do is to approximate the gradient of the log-likelihood  $\nabla_{\mathbf{x}_t} \log p(\mathbf{y}|\mathbf{x}_t)$ .
- However, the likelihood is given by an integral

$$p(\mathbf{y}|\mathbf{x}_t) = \int p(\mathbf{y}|\mathbf{x}_0) p_0(\mathbf{x}_0|\mathbf{x}_t) d\mathbf{x}_0.$$

This is hard to compute, and it can be shown that doing so is intractable [GJP<sup>+</sup>24]. So, we have to do something else.

### 3 Taxonomy of Methods

- The survey paper includes about 30 methods or so.
- It attempts to categorize these according to three sets of criteria.
  - **What the method does.** This is the main criterion. There are 5 categories.
    1. **Explicit approximation for likelihood term.** Approximate  $\nabla \log p(\mathbf{y}|\mathbf{x}_t)$  with a closed-form expression.
    2. **Variational inference.** Approximate the posterior  $p(\mathbf{x}|\mathbf{y})$  with a simpler, tractable, parameterized distribution. The parameters of these distribution must then be optimized.
    3. **CSGM-type methods.** Optimize the noise  $\mathbf{x}_T$  used to start the diffusion process.
    4. **Asymptotically exact methods.** Try to sample from the true posterior distribution  $p(\mathbf{x}|\mathbf{y})$  by either using Markov chain Monte Carlo (MCMC) or sequential Monte Carlo (SMC).
    5. **Others.** This category houses methods that do not fall into the previous categories.
  - **Optimization techniques used.**
    1. **Grad.** Update  $\mathbf{x}_t$  with a gradient.
    2. **Proj.** Project  $\mathbf{x}_t$  or  $E[\mathbf{x}_0|\mathbf{x}_t]$  to the “measurement subspace.”
    3. **Samp.** Sample the next particle by defining a proposal distribution and propagate multiple chains of particles.
    4. **Opt.** Define and solve an optimization problem at every step of the sampling process. Alternatively, define a global optimization problem that encompasses all timesteps and then solve it.

A method can use more than one optimization techniques and so can belong to multiple categories.

- **Type of inverse problem the method can solve.**
  - \* An inverse problem can be linear or non-linear based on whether the operator  $\mathcal{A}$  is linear or not.
  - \* An inverse problem may not have noise ( $\sigma_{\mathbf{y}} = 0$ ) or have it ( $\sigma_{\mathbf{y}} > 0$ ).
  - \* An inverse problem can be blind or non-blind based on whether information on  $\mathcal{A}$  is available or not. In a typically blind problems, the type of operator  $\mathcal{A}$  is known, but its parameters are unknown.

## 4 Explicit Approximation to the Likelihood Term

- In general, the approximation has the following form

$$\nabla_{\mathbf{x}_t} \log p(\mathbf{y}|\mathbf{x}_t) \approx -\frac{\mathcal{L}_t \mathcal{M}_t}{\mathcal{G}_t}.$$

Here,

- $\mathcal{M}_t \in \mathbb{R}^m$  represents an error vector that measures the discrepancy between the real measurement  $\mathbf{y}$  and the measurement restored from  $\mathbf{x}_t$ .
- $\mathcal{L}_t \in \mathbb{R}^{n \times m}$  is a matrix the projects the error vector  $\mathcal{M}_t \in \mathbb{R}^m$  back into  $\mathbb{R}^n$ .
- $\mathcal{G}_t$  is a scalar scaling factor.

### 4.1 Score-Based Annealed Langevin Dynamics (Score ALD)

- Proposed by Jalal et al. [JAD<sup>+</sup>21].
- Works only on linear problems with noise.
  - Let  $\mathcal{A}(\mathbf{x}) = A\mathbf{x}$  where  $A \in \mathbb{R}^{m \times n}$ .
- Makes the following approximation:

$$\nabla_{\mathbf{x}_t} \log p(\mathbf{y}|\mathbf{x}_t) \approx -\frac{A^T(\mathbf{y} - A\mathbf{x}_t)}{\sigma_{\mathbf{y}}^2 + \gamma_t^2}$$

where  $\gamma_t$  is a parameter to be tuned.

- The method guides the sample in the opposite direction of the measurement error  $\mathcal{M}_t = \mathbf{y} - A\mathbf{x}_t$  after being projected back with  $\mathcal{L}_t = A^T$ . So, this method is of category Proj.

### 4.2 Diffusion Posterior Sampling (DPS)

- Proposed by Chung et al. [CKM<sup>+</sup>24].
- DPS makes the following assumption.

$$\nabla_{\mathbf{x}_t} \log p(\mathbf{y}|\mathbf{X}_t = \mathbf{x}_t) \approx \nabla_{\mathbf{x}_t} \log p(\mathbf{y}|\mathbf{X}_0 = E[\mathbf{X}_0|\mathbf{X}_t = \mathbf{x}_t]).$$

- Because

$$p(\mathbf{y}|\mathbf{X}_0 = E[\mathbf{X}_0|\mathbf{X}_t = \mathbf{x}_t]) = \mathcal{N}(\mathbf{y}; \mathcal{A}(E[\mathbf{X}_0|\mathbf{X}_t = \mathbf{x}_t]), \sigma_{\mathbf{y}}^2 I),$$

we have that

$$\begin{aligned} \nabla_{\mathbf{x}_t} \log p(\mathbf{y}|\mathbf{X}_t = \mathbf{x}_t) &= \nabla_{\mathbf{x}_t} \log \mathcal{N}(\mathbf{y}; \mathcal{A}(E[\mathbf{X}_0|\mathbf{X}_t = \mathbf{x}_t]), \sigma_{\mathbf{y}}^2 I) \\ &= \nabla_{\mathbf{x}_t} \left( \frac{1}{2\sigma_{\mathbf{y}}^2} \|\mathbf{y} - \mathcal{A}(E[\mathbf{X}_0|\mathbf{X}_t = \mathbf{x}_t])\|^2 \right) \\ &= -\frac{1}{2\sigma_{\mathbf{y}}^2} (\nabla_{\mathbf{x}_t} \mathcal{A}(E[\mathbf{X}_0|\mathbf{X}_t = \mathbf{x}_t]))^T (\mathbf{y} - \mathcal{A}(E[\mathbf{X}_0|\mathbf{X}_t = \mathbf{x}_t])). \end{aligned}$$

So,

$$- \mathcal{G}_t = 1/(2\sigma_{\mathbf{y}}^2),$$



- $\mathcal{L}_t = (\nabla_{\mathbf{x}_t} \mathcal{A}(E[\mathbf{X}_0|\mathbf{X}_t = \mathbf{x}_t]))^T$ , and
- $\mathcal{M}_t = \mathbf{y} - \mathcal{A}(E[\mathbf{X}_0|\mathbf{X}_t = \mathbf{x}_t])$ .

- In practice, DPS does not use  $\mathcal{G}_t = 1/(2\sigma_y^2)$ . Instead, it proposes using scaling factor

$$\mathcal{G}_t = \zeta_t = \frac{\zeta'}{\|\mathbf{y} - \mathcal{A}(E[\mathbf{X}_0|\mathbf{X}_t = \mathbf{x}_t])\|}$$

where  $\zeta'$  is a constant.

## References

- [And82] B. D. Anderson, *Reverse-time diffusion equation models*, Stochastic Processes and their Applications **12** (1982), no. 3, 313–326.
- [BM18] Y. Blau and T. Michaeli, *The perception-distortion tradeoff*, 2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition, IEEE, June 2018, p. 6228–6237.
- [CKM<sup>+</sup>24] H. Chung, J. Kim, M. T. Mccann, M. L. Klasky, and J. C. Ye, *Diffusion posterior sampling for general noisy inverse problems*, 2024, [arXiv:2209.14687](#) [stat.ML].
- [DCL<sup>+</sup>24] G. Daras, H. Chung, C.-H. Lai, Y. Mitsufuji, J. C. Ye, P. Milanfar, A. G. Dimakis, and M. Delbracio, *A survey on diffusion models for inverse problems*, 2024, [arXiv:2410.00083](#) [cs.LG].
- [DM24] M. Delbracio and P. Milanfar, *Inversion by direct iteration: An alternative to denoising diffusion for image restoration*, 2024, [arXiv:2303.11435](#) [eess.IV].
- [GJP<sup>+</sup>24] S. Gupta, A. Jalal, A. Parulekar, E. Price, and Z. Xun, *Diffusion posterior sampling is computationally intractable*, 2024, [arXiv:2402.12727](#) [cs.LG].
- [JAD<sup>+</sup>21] A. Jalal, M. Arvinte, G. Daras, E. Price, A. G. Dimakis, and J. I. Tamir, *Robust compressed sensing mri with deep generative priors*, 2021, [arXiv:2108.01368](#) [cs.LG].
- [KAAL22] T. Karras, M. Aittala, T. Aila, and S. Laine, *Elucidating the design space of diffusion-based generative models*, 2022, [arXiv:2206.00364](#) [cs.CV].
- [LGZ<sup>+</sup>23a] Z. Luo, F. K. Gustafsson, Z. Zhao, J. Sjölund, and T. B. Schön, *Image restoration with mean-reverting stochastic differential equations*, 2023, [arXiv:2301.11699](#) [cs.LG].
- [LGZ<sup>+</sup>23b] ———, *Refusion: Enabling large-size realistic image restoration with latent-space diffusion models*, 2023, [arXiv:2304.08291](#) [cs.CV].
- [LYC<sup>+</sup>21] H. Li, Y. Yang, M. Chang, H. Feng, Z. Xu, Q. Li, and Y. Chen, *Srdiff: Single image super-resolution with diffusion probabilistic models*, 2021, [arXiv:2104.14951](#) [cs.CV].
- [SCC<sup>+</sup>22] C. Saharia, W. Chan, H. Chang, C. A. Lee, J. Ho, T. Salimans, D. J. Fleet, and M. Norouzi, *Palette: Image-to-image diffusion models*, 2022, [arXiv:2111.05826](#) [cs.CV].
- [SHC<sup>+</sup>21] C. Saharia, J. Ho, W. Chan, T. Salimans, D. J. Fleet, and M. Norouzi, *Image super-resolution via iterative refinement*, [arXiv:2104.07636](#) (2021).
- [SME22] J. Song, C. Meng, and S. Ermon, *Denoising diffusion implicit models*, 2022, [arXiv:2010.02502](#) [cs.LG].
- [SSDK<sup>+</sup>21] Y. Song, J. Sohl-Dickstein, D. P. Kingma, A. Kumar, S. Ermon, and B. Poole, *Score-based generative modeling through stochastic differential equations*, 2021, [arXiv:2011.13456](#) [cs.LG].

- [WDT<sup>+</sup>21] J. Whang, M. Delbracio, H. Talebi, C. Saharia, A. G. Dimakis, and P. Milanfar, *Deblurring via stochastic refinement*, 2021, [arXiv:2112.02475](#) [cs.CV].
- [ZRA23] L. Zhang, A. Rao, and M. Agrawala, *Adding conditional control to text-to-image diffusion models*, 2023, [arXiv:2302.05543](#) [cs.CV].