

# Flow Matching for Generative Modeling

Pramook Khungurn

July 22, 2024

This note was written as I read the “Flow Matching for Generative Modeling” paper by Lipman et al. [5].

## 1 Background

- A data item is denoted by  $x = (x^1, x^2, \dots, x^d) \in \mathbb{R}^d$ .
- A **probability density path** is a function  $p : [0, 1] \times \mathbb{R}^d \rightarrow \mathbb{R}^+ \cup \{0\}$  such that each  $p(t, \cdot)$  is a probability density function on  $\mathbb{R}^d$ . In other words, it holds that

$$\int p(t, x) dx = 1$$

for all  $t \in [0, 1]$ .

- For a time dependent function  $f : [0, 1] \times \mathbb{R}^d \rightarrow R$  for some range set  $R$ , we may write  $f(t, x)$  as  $f_t(x)$  to emphasize time dependence. Moreover, we can refer to  $f_t : \mathbb{R}^d \rightarrow R$  as a function in its own right.
  - With this, we may say that  $p_t$  is a probability distribution on  $\mathbb{R}^d$ .
- A **time-dependent vector field** is a function  $v : [0, 1] \times \mathbb{R}^d \rightarrow \mathbb{R}^d$ .
- Given a time dependent vector field  $v$ , its **flow** is another vector field  $\phi : [0, 1] \times \mathbb{R}^d \rightarrow \mathbb{R}^d$  defined by the ordinary differential equation

$$\frac{d}{dt} \phi_t(x) = v_t(\phi_t(x))$$

and the initial condition  $\phi_0(x) = x$ .

- In other words,  $\phi_t(x)$  is the position at time  $t$  of the particle that starts at  $x$  at time 0 and follows the trajectory defined by taking  $v$  as the time-dependent vector field.
- Chen et al. proposed the **neural ordinary differential equation** model [1]. The idea is to model the vector field  $v$  with a neural network  $v_t(x; \theta)$ . We then train it so that  $\phi_1$  has the property that we want.
  - If you want a refresher on neural ODE, then read my previous note on it [2].
- A neural ODE can be used to transform a probability distribution to another. Say, we start with a probability distribution  $p_0$  on  $\mathbb{R}^d$ . Then, we do the following.
  - Sample  $x \sim p_0$ .
  - Compute  $x' = \phi_t(x)$  by integrating the neural ODE from 0 up to  $t$ .

Let us denote the probability density of  $x'$  by  $p_t$ . It follows that

$$p_t(x') = p_0(\phi_t^{-1}(x')) \det \left[ \frac{\partial \phi_t^{-1}}{\partial x}(x') \right]. \quad (1)$$

This is the standard formula for transformation of probability distribution. You can find this in section 3.1 on my notes on the subject [4].

- The formula in Equation (1) is not that great because there is an issue with variable capture. The  $x$  in  $\partial x$  is not a variable but a shorthand the positional argument of a function. I previously have introduced a system to deal with this kind of problem [3]. So, let's write the equation using that notation.

First, we note that  $\phi_t(x) = q(t, x)$  is a function that maps a  $(d+1)$ -dimensional space to a  $d$ -dimensional space. So, we can treat it in the same way as a function of signature  $\mathbb{R}^{d+1} \rightarrow \mathbb{R}^d$ . In other words, we can say that  $\phi$  takes  $d+1$  inputs. We can then divide the  $d+1$  inputs into two blocks.

- The first block is the first argument alone. Using Python slice notation, it is “1 : 2.” Using my “chapter” notation, it can be abbreviated as §1.
- The second block is the rest of the arguments. Using Python slice notation, it is “2 :  $d+2$ .” Using my “chapter” notation, it can be abbreviate as §2.

Hence, using my notation for partial derivatives, we can rewrite the equation as:

$$p_t(x') = p_0(\phi_t^{-1}(x')) \det \nabla_{\S 2} \phi_t^{-1}(x')$$

or, to be even briefer

$$p_t(x') = p_0(\phi_t^{-1}(x')) |\nabla_{\S 2} \phi_t^{-1}(x')|$$

- Let  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  and let  $v : \mathbb{R}^d \rightarrow \mathbb{R}^d$ . A **push-forward** (or a change of variable) of  $f$  according to  $v$  is a function of  $g : \mathbb{R}^d \rightarrow \mathbb{R}$  defined by

$$g(y) = f(v^{-1}(y)) |\nabla v^{-1}(y)|.$$

Here,  $\nabla$  denotes the derivative operator, which gets you the Jacobian matrix. We denote the push-forward of  $f$  according to  $v$  as  $[v]_* f$ .

- In the context of the discussion so far, we have that  $p_t = [\phi_t]_* p_0$ .
- When we use a neural ODE to transform a probability distribution from one to another (i.e., transforming  $p_0$  from  $p_1$ ), we call the resulting model a **continuous normalizing flow** model.

## 2 Flow Matching

### 2.1 Flow Matching Objective

- We want to use the above framework to transform a simple noise distribution  $p_0 = p_{\text{noise}}$  to a data distribution  $p_1 = p_{\text{data}}$ .
  - $p_{\text{noise}}$  is typically a Gaussian distribution  $p_0 = \mathcal{N}(0, I)$ .
  - As in most ML settings, we do not have access to the density function  $p_{\text{data}}$ , but we only have samples from the distribution.
- Suppose we know a probability path  $p_t$  and a time-dependent vector field  $u_t$  that has the following property:

- $p_0$  is the desired noise distribution, and  $p_1$  is the desired data distribution.
- $u_t$  is the vector field such that  $p_t = [u_t]_* p_0$ .

Suppose again that we want to model  $u_t$  with a neural network  $v_t(x; \theta)$ . Then, we may do it by minimizing the **flow matching objective**:

$$\mathcal{L}_{\text{FM}}(\theta) = E_{t \sim \mathcal{U}([0,1]), x \sim p_t} [\|u_t(x) - v_t(x; \theta)\|^2].$$

- The flow matching objective is usable if we know  $p_t$  and  $u_t$  before hand. However, in our settings, we do not know anything about  $u_t$ , and we only know  $p_0 = p_{\text{noise}}$  and  $p_1 = p_{\text{data}}$  but nothing in between.

## 2.2 Rewriting conditional paths and vector fields

- We still do not know what  $p_t$  exactly is, but let us engage in wishful thinking and try to dictate its form.
- Let  $x_1 \sim p_{\text{data}}$  be a data item. We look at the conditional probability density  $p_t(x|x_1)$ . Let us require that
  1.  $p_0(x|x_1) = p_{\text{noise}}(x)$ , and
  2.  $p_1(x|x_1) = \mathcal{N}(x; x_1, \sigma^2 I)$  where  $\sigma$  is a small number.
- Now, we have that

$$p_t(x) = \int p_t(x|x_1) p_{\text{data}}(x_1) dx_1.$$

Moreover, if we choose  $\sigma$  to be small enough, we would have that

$$p_1(x) \approx p_{\text{data}}(x).$$

•

## References

- [1] CHEN, R. T. Q., RUBANOVA, Y., BETTENCOURT, J., AND DUVENAUD, D. Neural ordinary differential equations, 2019.
- [2] KHUNGURN, P. Neural ordinary differential equations. <https://pkhungurn.github.io/notes/notes/ml/neural-ode/neural-ode.pdf>. Accessed: 2024-07-19.
- [3] KHUNGURN, P. Notation for multivariable derivatives. <https://pkhungurn.github.io/notes/notes/math/multivar-deriv-notations/multivar-deriv-notations.pdf>. Accessed: 2024-07-22.
- [4] KHUNGURN, P. Probability under transformation. <https://pkhungurn.github.io/notes/notes/gfx/pdf-transform/pdf-transform.pdf>. Accessed: 2024-07-22.
- [5] LIPMAN, Y., CHEN, R. T. Q., BEN-HAMU, H., NICKEL, M., AND LE, M. Flow matching for generative modeling, 2023.